

Keras -- MLPs on MNIST

In [0]:

```
# Checking if GPU is available
from tensorflow.python.client import device_lib
print(device_lib.list_local_devices())
```

```
[name: "/device:CPU:0"
device_type: "CPU"
memory_limit: 268435456
locality {
}
incarnation: 12961759298813835808
, name: "/device:XLA_CPU:0"
device_type: "XLA_CPU"
memory_limit: 17179869184
locality {
}
incarnation: 7259935812420886672
physical_device_desc: "device: XLA_CPU device"
, name: "/device:XLA_GPU:0"
device_type: "XLA_GPU"
memory_limit: 17179869184
locality {
}
incarnation: 10587928147365517172
physical_device_desc: "device: XLA_GPU device"
, name: "/device:GPU:0"
device_type: "GPU"
memory_limit: 13036516148
locality {
  bus_id: 1
  links {
  }
}
incarnation: 5943975532915389494
physical_device_desc: "device: 0, name: Tesla T4, pci bus id: 0000:00:04.0, compute capability: 7.5"
]
```

In [0]:

```
# Mounting Google drive to save our trained models
from google.colab import drive
drive.mount('/content/gdrive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdqf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Ab&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code

Enter your authorization code:
.....

Mounted at /content/gdrive



In [0]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal # weight initializer
```

Using TensorFlow backend.

In [0]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [0]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [0]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [0]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [0]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  3  18  18  18  126  136  175  26  166  255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94  154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
 0  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0]
```


0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.

In [0]:

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

In [0]:

```
# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)
```

```
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [0]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1] # passed to the first hidden layer specifying the number of inputs
it should expect

batch_size = 128
nb_epoch = 20
```

Input -> Softmax:

In [0]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

```
WARNING: Logging before flag parsing goes to stderr.
W0804 16:31:32.464691 139651759564672 deprecation_wrapper.py:119] From
/usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:74: The name
tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

W0804 16:31:32.470091 139651759564672 deprecation_wrapper.py:119] From
/usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:517: The name
tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

W0804 16:31:32.474256 139651759564672 deprecation_wrapper.py:119] From
/usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4138: The name
tf.random_uniform is deprecated. Please use tf.random.uniform instead.
```

In [0]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, step
s_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
W0804 16:31:32.505516 139651759564672 deprecation_wrapper.py:119] From
/usr/local/lib/python3.6/dist-packages/keras/optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
W0804 16:31:32.528835 139651759564672 deprecation_wrapper.py:119] From
/usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3295: The name tf.log is deprecated. Please use tf.math.log instead.
```

```
W0804 16:31:32.617900 139651759564672 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
Instructions for updating:
```

```
Use tf.where in 2.0, which has the same broadcast rule as np.where
```

```
W0804 16:31:32.658567 139651759564672 deprecation_wrapper.py:119] From
/usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:986: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 37us/step - loss: 1.2731 - acc: 0.7043 - val_loss: 0.8063 - val_acc: 0.8367

Epoch 2/20

60000/60000 [=====] - 1s 23us/step - loss: 0.7110 - acc: 0.8448 - val_loss: 0.6050 - val_acc: 0.8658

Epoch 3/20

60000/60000 [=====] - 1s 23us/step - loss: 0.5838 - acc: 0.8609 - val_loss: 0.5247 - val_acc: 0.8742

Epoch 4/20

60000/60000 [=====] - 1s 22us/step - loss: 0.5231 - acc: 0.8698 - val_loss: 0.4796 - val_acc: 0.8805

```

Epoch 5/20
60000/60000 [=====] - 1s 22us/step - loss: 0.4860 - acc: 0.8757 -
val_loss: 0.4499 - val_acc: 0.8858
Epoch 6/20
60000/60000 [=====] - 1s 23us/step - loss: 0.4606 - acc: 0.8803 -
val_loss: 0.4292 - val_acc: 0.8894
Epoch 7/20
60000/60000 [=====] - 1s 22us/step - loss: 0.4418 - acc: 0.8844 -
val_loss: 0.4128 - val_acc: 0.8925
Epoch 8/20
60000/60000 [=====] - 1s 23us/step - loss: 0.4270 - acc: 0.8866 -
val_loss: 0.3999 - val_acc: 0.8942
Epoch 9/20
60000/60000 [=====] - 1s 22us/step - loss: 0.4151 - acc: 0.8889 -
val_loss: 0.3898 - val_acc: 0.8954
Epoch 10/20
60000/60000 [=====] - 1s 23us/step - loss: 0.4053 - acc: 0.8910 -
val_loss: 0.3813 - val_acc: 0.8979
Epoch 11/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3969 - acc: 0.8930 -
val_loss: 0.3741 - val_acc: 0.8991
Epoch 12/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3897 - acc: 0.8940 -
val_loss: 0.3677 - val_acc: 0.9013
Epoch 13/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3834 - acc: 0.8957 -
val_loss: 0.3618 - val_acc: 0.9026
Epoch 14/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3778 - acc: 0.8968 -
val_loss: 0.3573 - val_acc: 0.9033
Epoch 15/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3728 - acc: 0.8980 -
val_loss: 0.3526 - val_acc: 0.9045
Epoch 16/20
60000/60000 [=====] - 1s 22us/step - loss: 0.3683 - acc: 0.8986 -
val_loss: 0.3487 - val_acc: 0.9052
Epoch 17/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3643 - acc: 0.9000 -
val_loss: 0.3451 - val_acc: 0.9061
Epoch 18/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3606 - acc: 0.9007 -
val_loss: 0.3419 - val_acc: 0.9075
Epoch 19/20
60000/60000 [=====] - 1s 22us/step - loss: 0.3572 - acc: 0.9016 -
val_loss: 0.3387 - val_acc: 0.9079
Epoch 20/20
60000/60000 [=====] - 1s 22us/step - loss: 0.3540 - acc: 0.9023 -
val_loss: 0.3362 - val_acc: 0.9087

```

Note: The `keras.evaluate()` function will give you the loss value for every batch. The `keras.predict()` function will give you the actual predictions for all samples in a batch, for all batches. So even if you use the same data, the differences will be there because the value of a loss function will be almost always different than the predicted values. These are two different things.

In [0]:

```

%matplotlib inline
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


```



```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.3362268351316452
Test accuracy: 0.9087



MLP + Sigmoid activation + SGDOptimizer

In [0]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [0]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 30us/step - loss: 2.2726 - acc: 0.2015 - val_loss: 2.2266 - val_acc: 0.3424
Epoch 2/20
60000/60000 [=====] - 2s 26us/step - loss: 2.1825 - acc: 0.4357 - val_loss: 2.1259 - val_acc: 0.4958
Epoch 3/20
60000/60000 [=====] - 2s 26us/step - loss: 2.0662 - acc: 0.5613 -

```

val_loss: 1.9858 - val_acc: 0.6221
Epoch 4/20
60000/60000 [=====] - 2s 26us/step - loss: 1.9039 - acc: 0.6235 -
val_loss: 1.7950 - val_acc: 0.6902
Epoch 5/20
60000/60000 [=====] - 2s 26us/step - loss: 1.6967 - acc: 0.6728 -
val_loss: 1.5688 - val_acc: 0.7118
Epoch 6/20
60000/60000 [=====] - 2s 26us/step - loss: 1.4701 - acc: 0.7121 -
val_loss: 1.3422 - val_acc: 0.7296
Epoch 7/20
60000/60000 [=====] - 2s 26us/step - loss: 1.2590 - acc: 0.7438 -
val_loss: 1.1481 - val_acc: 0.7611
Epoch 8/20
60000/60000 [=====] - 2s 26us/step - loss: 1.0857 - acc: 0.7720 -
val_loss: 0.9954 - val_acc: 0.7788
Epoch 9/20
60000/60000 [=====] - 2s 26us/step - loss: 0.9530 - acc: 0.7917 -
val_loss: 0.8798 - val_acc: 0.8072
Epoch 10/20
60000/60000 [=====] - 2s 26us/step - loss: 0.8522 - acc: 0.8081 -
val_loss: 0.7928 - val_acc: 0.8204
Epoch 11/20
60000/60000 [=====] - 2s 26us/step - loss: 0.7743 - acc: 0.8209 -
val_loss: 0.7243 - val_acc: 0.8320
Epoch 12/20
60000/60000 [=====] - 2s 27us/step - loss: 0.7129 - acc: 0.8311 -
val_loss: 0.6695 - val_acc: 0.8420
Epoch 13/20
60000/60000 [=====] - 2s 26us/step - loss: 0.6633 - acc: 0.8391 -
val_loss: 0.6249 - val_acc: 0.8488
Epoch 14/20
60000/60000 [=====] - 2s 26us/step - loss: 0.6227 - acc: 0.8463 -
val_loss: 0.5885 - val_acc: 0.8545
Epoch 15/20
60000/60000 [=====] - 2s 26us/step - loss: 0.5889 - acc: 0.8528 -
val_loss: 0.5580 - val_acc: 0.8589
Epoch 16/20
60000/60000 [=====] - 2s 26us/step - loss: 0.5605 - acc: 0.8577 -
val_loss: 0.5314 - val_acc: 0.8648
Epoch 17/20
60000/60000 [=====] - 2s 27us/step - loss: 0.5363 - acc: 0.8617 -
val_loss: 0.5092 - val_acc: 0.8691
Epoch 18/20
60000/60000 [=====] - 2s 26us/step - loss: 0.5156 - acc: 0.8661 -
val_loss: 0.4897 - val_acc: 0.8727
Epoch 19/20
60000/60000 [=====] - 2s 26us/step - loss: 0.4975 - acc: 0.8702 -
val_loss: 0.4727 - val_acc: 0.8758
Epoch 20/20
60000/60000 [=====] - 2s 27us/step - loss: 0.4816 - acc: 0.8731 -
val_loss: 0.4583 - val_acc: 0.8792

```

In [0]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lida
tion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy


```

```
# for each key in history.history we will have a list of length equal to number of epochs
```

```
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```


Test score: 0.4583240813970566


Test accuracy: 0.8792



In [0]:

```
%matplotlib inline  
w_after = model_sigmoid.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
out_w = w_after[4].flatten().reshape(-1,1)  
  
fig = plt.figure(figsize=(12,8))  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 3, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w,color='b')  
plt.xlabel('Hidden Layer 1')  
  
plt.subplot(1, 3, 2)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h2_w, color='r')  
plt.xlabel('Hidden Layer 2 ')  
  
plt.subplot(1, 3, 3)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=out_w,color='y')  
plt.xlabel('Output Layer ')  
  
plt.show()
```





MLP + Sigmoid activation + ADAM

In [0]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 35us/step - loss: 0.5409 - acc: 0.8575 -
val_loss: 0.2526 - val_acc: 0.9252
Epoch 2/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2222 - acc: 0.9349 -
val_loss: 0.1831 - val_acc: 0.9450
Epoch 3/20
60000/60000 [=====] - 2s 31us/step - loss: 0.1637 - acc: 0.9521 -
val_loss: 0.1407 - val_acc: 0.9577
Epoch 4/20
60000/60000 [=====] - 2s 31us/step - loss: 0.1252 - acc: 0.9631 -
val_loss: 0.1212 - val_acc: 0.9620
Epoch 5/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0978 - acc: 0.9711 -
val_loss: 0.1023 - val_acc: 0.9696
Epoch 6/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0779 - acc: 0.9772 -
val_loss: 0.0889 - val_acc: 0.9714
Epoch 7/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0636 - acc: 0.9814 -
val_loss: 0.0851 - val_acc: 0.9740
Epoch 8/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0510 - acc: 0.9853 -
val_loss: 0.0723 - val_acc: 0.9776
Epoch 9/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0422 - acc: 0.9874 -
val_loss: 0.0684 - val_acc: 0.9790
Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0331 - acc: 0.9908 -
val_loss: 0.0704 - val_acc: 0.9787
Epoch 11/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0282 - acc: 0.9919 -
val_loss: 0.0680 - val_acc: 0.9791
Epoch 12/20
```

```

60000/60000 [=====] - 2s 31us/step - loss: 0.0212 - acc: 0.9945 -
val_loss: 0.0616 - val_acc: 0.9809
Epoch 13/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0170 - acc: 0.9957 -
val_loss: 0.0629 - val_acc: 0.9807
Epoch 14/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0141 - acc: 0.9968 -
val_loss: 0.0689 - val_acc: 0.9791
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0118 - acc: 0.9973 -
val_loss: 0.0709 - val_acc: 0.9790
Epoch 16/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0087 - acc: 0.9981 -
val_loss: 0.0648 - val_acc: 0.9805
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0071 - acc: 0.9985 -
val_loss: 0.0705 - val_acc: 0.9812
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0057 - acc: 0.9989 -
val_loss: 0.0691 - val_acc: 0.9800
Epoch 19/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0050 - acc: 0.9990 -
val_loss: 0.0762 - val_acc: 0.9797
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0044 - acc: 0.9990 -
val_loss: 0.0737 - val_acc: 0.9801

```

In [0]:

```

%matplotlib inline
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0736692005243589

Test accuracy: 0.9801



In [0]:

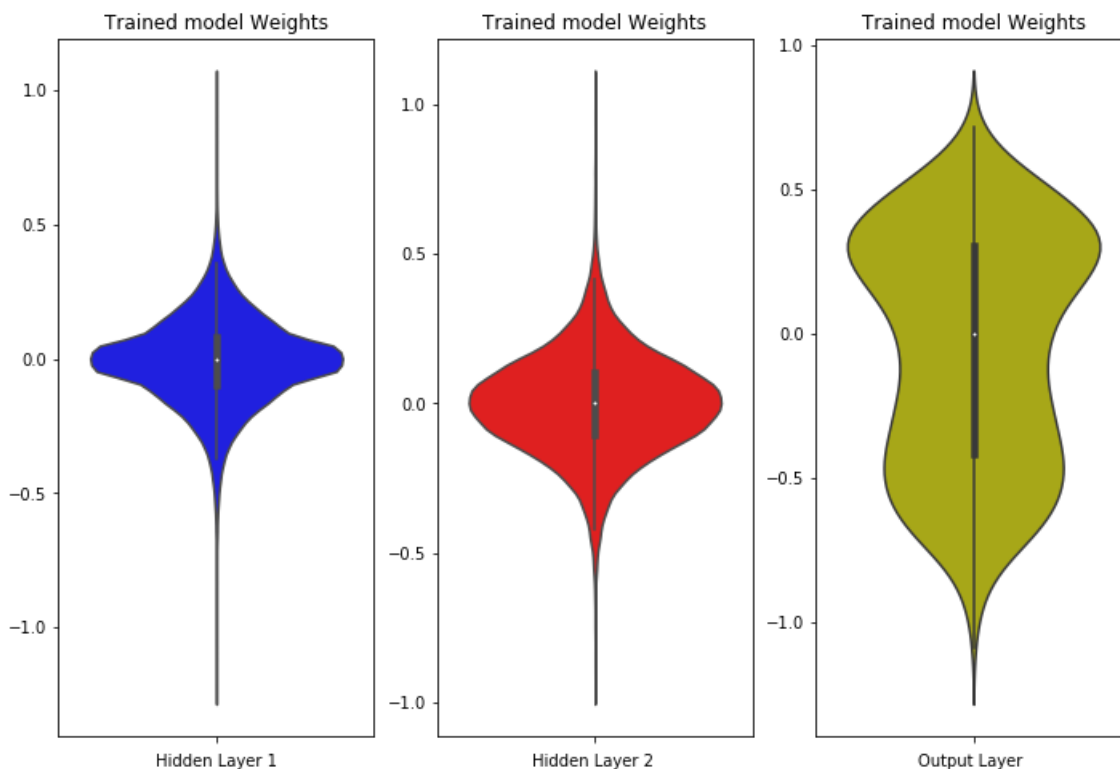
```
%matplotlib inline
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU +SGD

In [0]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0,\sigma) = N(0,0.062)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0,\sigma) = N(0,0.125)$ 
# out =>  $\sigma=\sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 
```

```

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [0]:

```

model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 2s 31us/step - loss: 0.7357 - acc: 0.7929 - val_loss: 0.3855 - val_acc: 0.8933
Epoch 2/20
60000/60000 [=====] - 2s 27us/step - loss: 0.3571 - acc: 0.8995 - val_loss: 0.2975 - val_acc: 0.9147
Epoch 3/20
60000/60000 [=====] - 2s 27us/step - loss: 0.2938 - acc: 0.9167 - val_loss: 0.2609 - val_acc: 0.9245
Epoch 4/20
60000/60000 [=====] - 2s 27us/step - loss: 0.2589 - acc: 0.9259 - val_loss: 0.2372 - val_acc: 0.9326
Epoch 5/20
60000/60000 [=====] - 2s 28us/step - loss: 0.2349 - acc: 0.9331 - val_loss: 0.2169 - val_acc: 0.9370
Epoch 6/20
60000/60000 [=====] - 2s 27us/step - loss: 0.2166 - acc: 0.9385 - val_loss: 0.2035 - val_acc: 0.9416
Epoch 7/20
60000/60000 [=====] - 2s 27us/step - loss: 0.2015 - acc: 0.9425 - val_loss: 0.1906 - val_acc: 0.9446
Epoch 8/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1886 - acc: 0.9467 - val_loss: 0.1815 - val_acc: 0.9475
Epoch 9/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1777 - acc: 0.9500 - val_loss: 0.1729 - val_acc: 0.9492
Epoch 10/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1682 - acc: 0.9525 - val_loss: 0.1659 - val_acc: 0.9507
Epoch 11/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1595 - acc: 0.9551 - val_loss: 0.1576 - val_acc: 0.9526
Epoch 12/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1519 - acc: 0.9569 - val_loss: 0.1513 - val_acc: 0.9540
Epoch 13/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1451 - acc: 0.9596 - val_loss: 0.1464 - val_acc: 0.9556
Epoch 14/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1389 - acc: 0.9611 - val_loss: 0.1418 - val_acc: 0.9582
Epoch 15/20

```

```

60000/60000 [=====] - 2s 27us/step - loss: 0.1330 - acc: 0.9631 -
val_loss: 0.1384 - val_acc: 0.9591
Epoch 16/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1277 - acc: 0.9646 -
val_loss: 0.1343 - val_acc: 0.9611
Epoch 17/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1228 - acc: 0.9660 -
val_loss: 0.1300 - val_acc: 0.9620
Epoch 18/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1180 - acc: 0.9671 -
val_loss: 0.1265 - val_acc: 0.9627
Epoch 19/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1140 - acc: 0.9685 -
val_loss: 0.1233 - val_acc: 0.9630
Epoch 20/20
60000/60000 [=====] - 2s 27us/step - loss: 0.1100 - acc: 0.9699 -
val_loss: 0.1210 - val_acc: 0.9644

```

In [0]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12101748173683882

Test accuracy: 0.9644



In [0]:

```

%matplotlib inline
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)

```



```


out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLU + ADAM

In [0]:

```

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920

dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

=====

Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 37us/step - loss: 0.2257 - acc: 0.9323 -
val_loss: 0.1141 - val_acc: 0.9634
Epoch 2/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0841 - acc: 0.9741 -
val_loss: 0.0790 - val_acc: 0.9759
Epoch 3/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0521 - acc: 0.9840 -
val_loss: 0.0830 - val_acc: 0.9751
Epoch 4/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0350 - acc: 0.9891 -
val_loss: 0.0743 - val_acc: 0.9775
Epoch 5/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0256 - acc: 0.9918 -
val_loss: 0.0639 - val_acc: 0.9804
Epoch 6/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0194 - acc: 0.9942 -
val_loss: 0.0672 - val_acc: 0.9800
Epoch 7/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0191 - acc: 0.9939 -
val_loss: 0.0847 - val_acc: 0.9746
Epoch 8/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0153 - acc: 0.9949 -
val_loss: 0.0646 - val_acc: 0.9810
Epoch 9/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0140 - acc: 0.9950 -
val_loss: 0.0900 - val_acc: 0.9754
Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0110 - acc: 0.9962 -
val_loss: 0.1109 - val_acc: 0.9720
Epoch 11/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0090 - acc: 0.9970 -
val_loss: 0.0810 - val_acc: 0.9805
Epoch 12/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0120 - acc: 0.9961 -
val_loss: 0.0868 - val_acc: 0.9787
Epoch 13/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0099 - acc: 0.9966 -
val_loss: 0.0769 - val_acc: 0.9809
Epoch 14/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0105 - acc: 0.9965 -
val_loss: 0.0938 - val_acc: 0.9789
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0073 - acc: 0.9977 -
val_loss: 0.0889 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0064 - acc: 0.9978 -
val_loss: 0.0902 - val_acc: 0.9802
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0072 - acc: 0.9979 -
val_loss: 0.0987 - val_acc: 0.9790
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0078 - acc: 0.9975 -
val_loss: 0.0963 - val_acc: 0.9785
Epoch 19/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0089 - acc: 0.9970 -
val_loss: 0.0985 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0054 - acc: 0.9983 -
val_loss: 0.0951 - val_acc: 0.9796

```

```
In [0]:
```

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])

```

```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09506532158235209

Test accuracy: 0.9796



In [0]:

```

%matplotlib inline
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

Trained model Weights

Trained model Weights

Trained model Weights



MLP + Batch-Norm on hidden Layers + AdamOptimizer

In [0]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		

Trainable params: 470,154
Non-trainable params: 1,280

In [0]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 62us/step - loss: 0.3042 - acc: 0.9093 -
val_loss: 0.2186 - val_acc: 0.9327
Epoch 2/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1785 - acc: 0.9474 -
val_loss: 0.1755 - val_acc: 0.9468
Epoch 3/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1385 - acc: 0.9595 -
val_loss: 0.1494 - val_acc: 0.9534
Epoch 4/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1130 - acc: 0.9668 -
val_loss: 0.1292 - val_acc: 0.9617
Epoch 5/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0956 - acc: 0.9711 -
val_loss: 0.1328 - val_acc: 0.9597
Epoch 6/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0817 - acc: 0.9752 -
val_loss: 0.1192 - val_acc: 0.9616
Epoch 7/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0704 - acc: 0.9784 -
val_loss: 0.1181 - val_acc: 0.9651
Epoch 8/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0604 - acc: 0.9808 -
val_loss: 0.1128 - val_acc: 0.9657
Epoch 9/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0498 - acc: 0.9843 -
val_loss: 0.1026 - val_acc: 0.9682
Epoch 10/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0435 - acc: 0.9866 -
val_loss: 0.1022 - val_acc: 0.9701
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0399 - acc: 0.9879 -
val_loss: 0.1038 - val_acc: 0.9692
Epoch 12/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0372 - acc: 0.9878 -
val_loss: 0.1052 - val_acc: 0.9691
Epoch 13/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0301 - acc: 0.9902 -
val_loss: 0.0957 - val_acc: 0.9712
Epoch 14/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0269 - acc: 0.9911 -
val_loss: 0.1068 - val_acc: 0.9714
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0271 - acc: 0.9908 -
val_loss: 0.0998 - val_acc: 0.9738
Epoch 16/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0237 - acc: 0.9923 -
val_loss: 0.0985 - val_acc: 0.9738
Epoch 17/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0225 - acc: 0.9931 -
val_loss: 0.0975 - val_acc: 0.9731
Epoch 18/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0205 - acc: 0.9931 -
val_loss: 0.0935 - val_acc: 0.9748
Epoch 19/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0177 - acc: 0.9943 -
val_loss: 0.0903 - val_acc: 0.9754
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0161 - acc: 0.9946 -
val_loss: 0.0952 - val_acc: 0.9749
```

In [0]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09516252700814512

Test accuracy: 0.9749



In [0]:

```

%matplotlib inline
w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



5. MLP + Dropout + AdamOptimizer

In [0]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

W0804 16:35:36.860419 139651759564672 deprecation.py:506] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664

dense_18 (Dense)	(None, 128)	16384
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
=====		
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		
=====		

In [0]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 67us/step - loss: 0.6652 - acc: 0.7950 -
val_loss: 0.2836 - val_acc: 0.9164
Epoch 2/20
60000/60000 [=====] - 3s 52us/step - loss: 0.4251 - acc: 0.8695 -
val_loss: 0.2539 - val_acc: 0.9236
Epoch 3/20
60000/60000 [=====] - 3s 52us/step - loss: 0.3772 - acc: 0.8852 -
val_loss: 0.2323 - val_acc: 0.9309
Epoch 4/20
60000/60000 [=====] - 3s 52us/step - loss: 0.3527 - acc: 0.8953 -
val_loss: 0.2187 - val_acc: 0.9365
Epoch 5/20
60000/60000 [=====] - 3s 52us/step - loss: 0.3349 - acc: 0.8995 -
val_loss: 0.2018 - val_acc: 0.9415
Epoch 6/20
60000/60000 [=====] - 3s 52us/step - loss: 0.3184 - acc: 0.9046 -
val_loss: 0.1969 - val_acc: 0.9428
Epoch 7/20
60000/60000 [=====] - 3s 52us/step - loss: 0.3002 - acc: 0.9099 -
val_loss: 0.1915 - val_acc: 0.9432
Epoch 8/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2929 - acc: 0.9111 -
val_loss: 0.1806 - val_acc: 0.9483
Epoch 9/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2798 - acc: 0.9172 -
val_loss: 0.1759 - val_acc: 0.9486
Epoch 10/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2677 - acc: 0.9191 -
val_loss: 0.1643 - val_acc: 0.9514
Epoch 11/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2589 - acc: 0.9218 -
val_loss: 0.1567 - val_acc: 0.9532
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2459 - acc: 0.9253 -
val_loss: 0.1485 - val_acc: 0.9546
Epoch 13/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2336 - acc: 0.9302 -
val_loss: 0.1470 - val_acc: 0.9568
Epoch 14/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2257 - acc: 0.9325 -
val_loss: 0.1392 - val_acc: 0.9589
Epoch 15/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2158 - acc: 0.9349 -
val_loss: 0.1334 - val_acc: 0.9602
Epoch 16/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2055 - acc: 0.9387 -
val_loss: 0.1287 - val_acc: 0.9627
Epoch 17/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1996 - acc: 0.9395 -
val_loss: 0.1232 - val_acc: 0.9631
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1891 - acc: 0.9429 -
val_loss: 0.1133 - val_acc: 0.9659
Epoch 19/20
```



```
60000/60000 [=====] - 3s 53us/step - loss: 0.1833 - acc: 0.9449 -
val_loss: 0.1158 - val_acc: 0.9664
Epoch 20/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1769 - acc: 0.9468 -
val_loss: 0.1058 - val_acc: 0.9683
```

In [0]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10577300073988735

Test accuracy: 0.9683



In [0]:

```
%matplotlib inline
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [0]:

```
from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [0]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
```

```
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter
```

```
grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:1978: FutureWarning: The
default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this
warning.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

In [0]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.975817 using {'activ': 'relu'}
0.975017 (0.001059) with: {'activ': 'sigmoid'}
0.975817 (0.002400) with: {'activ': 'relu'}
```

With 2 Hiden layers

Architecture used: 784 (input) - 256 - 64 - 10 (ouput)

MLP + ReLu + Adam

In [0]:

```
model2_relu = Sequential()

model2_relu.add(Dense(256, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_
uniform'))
model2_relu.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_u
niform'))
model2_relu.add(Dense(output_dim, activation = 'softmax'))

model2_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 256)	200960
dense_42 (Dense)	(None, 64)	16448
dense_43 (Dense)	(None, 10)	650
Total params: 218,058		
Trainable params: 218,058		
Non-trainable params: 0		

In [0]:

```
model2_relu.compile(optimizer = 'adam', metrics = ['accuracy'], loss = 'categorical_crossentropy')

history = model2_relu.fit(X_train, Y_train, batch_size= batch_size, epochs = nb_epoch, verbose = 1,
validation_data = (X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 3s 54us/step - loss: 0.2851 - acc: 0.9207 -
val_loss: 0.1351 - val_acc: 0.9590
Epoch 2/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1053 - acc: 0.9687 -
```

```

val_loss: 0.1022 - val_acc: 0.9695
Epoch 3/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0706 - acc: 0.9788 -
val_loss: 0.0691 - val_acc: 0.9781
Epoch 4/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0507 - acc: 0.9847 -
val_loss: 0.0719 - val_acc: 0.9769
Epoch 5/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0402 - acc: 0.9874 -
val_loss: 0.0722 - val_acc: 0.9777
Epoch 6/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0306 - acc: 0.9906 -
val_loss: 0.0621 - val_acc: 0.9820
Epoch 7/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0232 - acc: 0.9927 -
val_loss: 0.0688 - val_acc: 0.9798
Epoch 8/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0194 - acc: 0.9941 -
val_loss: 0.0767 - val_acc: 0.9781
Epoch 9/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0178 - acc: 0.9944 -
val_loss: 0.0693 - val_acc: 0.9807
Epoch 10/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0128 - acc: 0.9962 -
val_loss: 0.0662 - val_acc: 0.9807
Epoch 11/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0117 - acc: 0.9965 -
val_loss: 0.0684 - val_acc: 0.9820
Epoch 12/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0117 - acc: 0.9963 -
val_loss: 0.0728 - val_acc: 0.9801
Epoch 13/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0088 - acc: 0.9970 -
val_loss: 0.0742 - val_acc: 0.9818
Epoch 14/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0112 - acc: 0.9962 -
val_loss: 0.0885 - val_acc: 0.9780
Epoch 15/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0072 - acc: 0.9976 -
val_loss: 0.0847 - val_acc: 0.9807
Epoch 16/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0092 - acc: 0.9969 -
val_loss: 0.0773 - val_acc: 0.9828
Epoch 17/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0064 - acc: 0.9980 -
val_loss: 0.0869 - val_acc: 0.9811
Epoch 18/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0067 - acc: 0.9978 -
val_loss: 0.1025 - val_acc: 0.9769
Epoch 19/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0122 - acc: 0.9961 -
val_loss: 0.0925 - val_acc: 0.9795
Epoch 20/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0058 - acc: 0.9981 -
val_loss: 0.1004 - val_acc: 0.9788

```

In [0]:

```

score = model2_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
litation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10035725422161268

Test accuracy: 0.9788



In [0]:


```
%matplotlib inline
w_after = model2_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + BatchNorm on hidden layers + Adam

In [0]:

```
model2_relu_batchnorm = Sequential()

model2_relu_batchnorm.add(Dense(256, activation = 'relu', input_dim = input_dim, kernel_initializer=
'glorot_uniform'))
model2_relu_batchnorm.add(BatchNormalization())

model2_relu_batchnorm.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer=
'glorot_uniform'))
model2_relu_batchnorm.add(BatchNormalization())

model2_relu_batchnorm.add(Dense(output_dim, activation = 'softmax'))

model2_relu_batchnorm.summary()
```

Layer (type)	Output Shape	Param #
dense_44 (Dense)	(None, 256)	200960
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
dense_45 (Dense)	(None, 64)	16448
batch_normalization_6 (Batch Normalization)	(None, 64)	256
dense_46 (Dense)	(None, 10)	650
Total params: 219,338		
Trainable params: 218,698		
Non-trainable params: 640		

In [0]:

```
model2_relu_batchnorm.compile(optimizer='adam', metrics = ['accuracy'], loss =
'categorical_crossentropy')

history = model2_relu_batchnorm.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, v
erbose = 1, validation_data = (X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 79us/step - loss: 0.2128 - acc: 0.9386 -
val_loss: 0.1270 - val_acc: 0.9603
Epoch 2/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0863 - acc: 0.9741 -
val_loss: 0.0921 - val_acc: 0.9732
Epoch 3/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0560 - acc: 0.9828 -
val_loss: 0.0807 - val_acc: 0.9747
Epoch 4/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0408 - acc: 0.9875 -
val_loss: 0.0719 - val_acc: 0.9773
Epoch 5/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0305 - acc: 0.9907 -
val_loss: 0.0722 - val_acc: 0.9772
```

```

Epoch 6/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0248 - acc: 0.9919 -
val_loss: 0.0712 - val_acc: 0.9789
Epoch 7/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0210 - acc: 0.9932 -
val_loss: 0.0685 - val_acc: 0.9812
Epoch 8/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0192 - acc: 0.9940 -
val_loss: 0.0754 - val_acc: 0.9768
Epoch 9/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0177 - acc: 0.9942 -
val_loss: 0.0757 - val_acc: 0.9792
Epoch 10/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0140 - acc: 0.9957 -
val_loss: 0.0735 - val_acc: 0.9771
Epoch 11/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0123 - acc: 0.9962 -
val_loss: 0.0715 - val_acc: 0.9789
Epoch 12/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0145 - acc: 0.9951 -
val_loss: 0.0835 - val_acc: 0.9785
Epoch 13/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0121 - acc: 0.9958 -
val_loss: 0.0859 - val_acc: 0.9779
Epoch 14/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0109 - acc: 0.9963 -
val_loss: 0.0727 - val_acc: 0.9816
Epoch 15/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0096 - acc: 0.9968 -
val_loss: 0.0830 - val_acc: 0.9791
Epoch 16/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0108 - acc: 0.9964 -
val_loss: 0.0841 - val_acc: 0.9790
Epoch 17/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0088 - acc: 0.9972 -
val_loss: 0.0796 - val_acc: 0.9815
Epoch 18/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0076 - acc: 0.9974 -
val_loss: 0.0782 - val_acc: 0.9807
Epoch 19/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0071 - acc: 0.9979 -
val_loss: 0.0703 - val_acc: 0.9815
Epoch 20/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0079 - acc: 0.9972 -
val_loss: 0.0851 - val_acc: 0.9786

```

In [0]:

```

score = model2_relu_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08512964024284957

Test accuracy: 0.9786



In [0]:


```
%matplotlib inline
w_after = model2_relu_batchnorm.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize = (12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Observation

After applying normalization, we see that weights in the output layer are mostly between 0.0 to 0.5 and fewer between 0.5 to 1.0

Note: I am going to experiment here. First, I will place batch norm before dropout and later I will place dropout before batchnorm. After doing this, I will compare the accuracy of these models.

This is some interesting discussion:

- <https://stackoverflow.com/questions/39691902/ordering-of-batch-normalization-and-dropout>
- <https://github.com/cvjena/cnn-models/issues/3>

MLP + Batch Norm + Dropout + Adam

In [0]:

```
model2_batchnorm_drop = Sequential()

model2_batchnorm_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer='random_normal'))
model2_batchnorm_drop.add(BatchNormalization())
model2_batchnorm_drop.add(Dropout(0.5))

model2_batchnorm_drop.add(Dense(64, activation='relu', kernel_initializer='random_normal'))
model2_batchnorm_drop.add(BatchNormalization())
model2_batchnorm_drop.add(Dropout(0.5))

model2_batchnorm_drop.add(Dense(output_dim, activation='softmax'))

model2_batchnorm_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 256)	200960
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0
dense_48 (Dense)	(None, 64)	16448
batch_normalization_8 (Batch Normalization)	(None, 64)	256
dropout_4 (Dropout)	(None, 64)	0
dense_49 (Dense)	(None, 10)	650
Total params: 219,338		
Trainable params: 218,698		
Non-trainable params: 640		

In [0]:

```
model2_batchnorm_drop.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model2_batchnorm_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 83us/step - loss: 0.5546 - acc: 0.8328 - val_loss: 0.1816 - val_acc: 0.9433

Epoch 2/20

60000/60000 [=====] - 3s 57us/step - loss: 0.2673 - acc: 0.9218 - val_loss: 0.1269 - val_acc: 0.9590

```

val_loss: 0.1205    val_acc: 0.9390
Epoch 3/20
60000/60000 [=====] - 3s 58us/step - loss: 0.2081 - acc: 0.9389 -
val_loss: 0.1117 - val_acc: 0.9676
Epoch 4/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1789 - acc: 0.9477 -
val_loss: 0.1001 - val_acc: 0.9704
Epoch 5/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1661 - acc: 0.9516 -
val_loss: 0.0845 - val_acc: 0.9744
Epoch 6/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1508 - acc: 0.9558 -
val_loss: 0.0901 - val_acc: 0.9713
Epoch 7/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1371 - acc: 0.9600 -
val_loss: 0.0850 - val_acc: 0.9734
Epoch 8/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1287 - acc: 0.9622 -
val_loss: 0.0772 - val_acc: 0.9774
Epoch 9/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1233 - acc: 0.9630 -
val_loss: 0.0713 - val_acc: 0.9790
Epoch 10/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1172 - acc: 0.9642 -
val_loss: 0.0749 - val_acc: 0.9789
Epoch 11/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1113 - acc: 0.9670 -
val_loss: 0.0730 - val_acc: 0.9782
Epoch 12/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1068 - acc: 0.9678 -
val_loss: 0.0737 - val_acc: 0.9785
Epoch 13/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1017 - acc: 0.9695 -
val_loss: 0.0693 - val_acc: 0.9799
Epoch 14/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0986 - acc: 0.9704 -
val_loss: 0.0684 - val_acc: 0.9796
Epoch 15/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0980 - acc: 0.9700 -
val_loss: 0.0685 - val_acc: 0.9795
Epoch 16/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0949 - acc: 0.9715 -
val_loss: 0.0684 - val_acc: 0.9790
Epoch 17/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0869 - acc: 0.9728 -
val_loss: 0.0727 - val_acc: 0.9789
Epoch 18/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0908 - acc: 0.9719 -
val_loss: 0.0659 - val_acc: 0.9800
Epoch 19/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0835 - acc: 0.9743 -
val_loss: 0.0635 - val_acc: 0.9812
Epoch 20/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0830 - acc: 0.9753 -
val_loss: 0.0617 - val_acc: 0.9827

```

In [0]:

```

score = model2_batchnorm_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.061712297642615155
Test accuracy: 0.9827



In [0]:


```
%matplotlib inline
w_after = model2_batchnorm_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize= (12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + Dropout + Batch Norm +Adam

In [0]:

```
model2_drop_batchnorm = Sequential()

model2_drop_batchnorm.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer='random_normal'))
model2_drop_batchnorm.add(BatchNormalization())
model2_drop_batchnorm.add(Dropout(0.5))

model2_drop_batchnorm.add(Dense(64, activation='relu', kernel_initializer='random_normal'))
model2_drop_batchnorm.add(BatchNormalization())
model2_drop_batchnorm.add(Dropout(0.5))

model2_drop_batchnorm.add(Dense(output_dim, activation='softmax'))

model2_drop_batchnorm.summary()
```

Layer (type)	Output Shape	Param #
dense_50 (Dense)	(None, 256)	200960
batch_normalization_9 (Batch Normalization)	(None, 256)	1024
dropout_5 (Dropout)	(None, 256)	0
dense_51 (Dense)	(None, 64)	16448
batch_normalization_10 (Batch Normalization)	(None, 64)	256
dropout_6 (Dropout)	(None, 64)	0
dense_52 (Dense)	(None, 10)	650
Total params: 219,338		
Trainable params: 218,698		
Non-trainable params: 640		

In [0]:

```
model2_drop_batchnorm.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model2_drop_batchnorm.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 82us/step - loss: 0.5392 - acc: 0.8366 - val_loss: 0.1746 - val_acc: 0.9465
Epoch 2/20
60000/60000 [=====] - 3s 57us/step - loss: 0.2621 - acc: 0.9240 - val_loss: 0.1209 - val_acc: 0.9629
Epoch 3/20
60000/60000 [=====] - 3s 57us/step - loss: 0.2077 - acc: 0.9392 - val_loss: 0.1033 - val_acc: 0.9683
Epoch 4/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1771 - acc: 0.9488
```

```

60000/60000 [=====] - 3s 57us/step - loss: 0.1771 - acc: 0.9488 -
val_loss: 0.0959 - val_acc: 0.9713
Epoch 5/20
60000/60000 [=====] - 3s 58us/step - loss: 0.1600 - acc: 0.9529 -
val_loss: 0.0874 - val_acc: 0.9735
Epoch 6/20
60000/60000 [=====] - 3s 58us/step - loss: 0.1478 - acc: 0.9562 -
val_loss: 0.0842 - val_acc: 0.9728
Epoch 7/20
60000/60000 [=====] - 3s 58us/step - loss: 0.1356 - acc: 0.9595 -
val_loss: 0.0859 - val_acc: 0.9740
Epoch 8/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1315 - acc: 0.9610 -
val_loss: 0.0747 - val_acc: 0.9762
Epoch 9/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1193 - acc: 0.9645 -
val_loss: 0.0740 - val_acc: 0.9775
Epoch 10/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1159 - acc: 0.9663 -
val_loss: 0.0711 - val_acc: 0.9789
Epoch 11/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1123 - acc: 0.9666 -
val_loss: 0.0691 - val_acc: 0.9797
Epoch 12/20
60000/60000 [=====] - 3s 58us/step - loss: 0.1061 - acc: 0.9692 -
val_loss: 0.0717 - val_acc: 0.9798
Epoch 13/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1004 - acc: 0.9705 -
val_loss: 0.0700 - val_acc: 0.9808
Epoch 14/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0971 - acc: 0.9708 -
val_loss: 0.0668 - val_acc: 0.9797
Epoch 15/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0935 - acc: 0.9719 -
val_loss: 0.0707 - val_acc: 0.9802
Epoch 16/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0893 - acc: 0.9727 -
val_loss: 0.0676 - val_acc: 0.9805
Epoch 17/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0894 - acc: 0.9726 -
val_loss: 0.0673 - val_acc: 0.9801
Epoch 18/20
60000/60000 [=====] - 4s 58us/step - loss: 0.0856 - acc: 0.9745 -
val_loss: 0.0630 - val_acc: 0.9821
Epoch 19/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0860 - acc: 0.9749 -
val_loss: 0.0667 - val_acc: 0.9801
Epoch 20/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0824 - acc: 0.9750 -
val_loss: 0.0660 - val_acc: 0.9813

```

In [0]:

```

score = model2_drop_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = nistory.nistory['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06602386177739827
Test accuracy: 0.9813



In [0]:


```
%matplotlib inline
w_after = model2_drop_batchnorm.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize= (12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + Batch Norm + Dropout + Adam

Test score: 0.0703403844089189

Test accuracy: 0.9796

MLP + Dropout + Batch Norm +Adam

Test score: 0.06499798040131573

Test accuracy: 0.9817

In [0]:

```
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Architecture Type", "Test score", "Test accuracy"]

x.add_row(["MLP + Batch Norm + Dropout + Adam", 0.0703403844089189, 0.9796])
x.add_row(["MLP + Dropout + Batch Norm +Adam", 0.06499798040131573, 0.9817])

print(x)
```

Architecture Type	Test score	Test accuracy
MLP + Batch Norm + Dropout + Adam	0.0703403844089189	0.9796
MLP + Dropout + Batch Norm +Adam	0.06499798040131573	0.9817

Observation: There is a slight improvement in the test accuracy if we use dropout before batchnorm.

With 3 Hiden layers

Architecture used: 784 (input) - 128 - 64 - 32 - 10 (ouput)

MLP + ReLu + Adam

In [0]:

```
model3_relu = Sequential()

model3_relu.add(Dense(128, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_normal'))
model3_relu.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_normal'))
model3_relu.add(Dense(32, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_normal'))
model3_relu.add(Dense(output_dim, activation = 'softmax'))

model3_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_53 (Dense)	(None, 128)	100480
dense_54 (Dense)	(None, 64)	8256
dense_55 (Dense)	(None, 32)	2080
dense_56 (Dense)	(None, 10)	330

Total params: 111,146
Trainable params: 111,146
Non-trainable params: 0

In [0]:

```
model3_relu.compile(optimizer = 'adam', metrics = ['accuracy'], loss = 'categorical_crossentropy')  
  
history = model3_relu.fit(X_train, Y_train, batch_size= batch_size, epochs = nb_epoch, verbose = 1,  
validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.3343 - acc: 0.9039 -  
val_loss: 0.1558 - val_acc: 0.9510  
Epoch 2/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.1296 - acc: 0.9618 -  
val_loss: 0.1152 - val_acc: 0.9641  
Epoch 3/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.0880 - acc: 0.9738 -  
val_loss: 0.1043 - val_acc: 0.9666  
Epoch 4/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0691 - acc: 0.9785 -  
val_loss: 0.0949 - val_acc: 0.9707  
Epoch 5/20  
60000/60000 [=====] - 2s 40us/step - loss: 0.0523 - acc: 0.9841 -  
val_loss: 0.0811 - val_acc: 0.9758  
Epoch 6/20  
60000/60000 [=====] - 2s 40us/step - loss: 0.0440 - acc: 0.9861 -  
val_loss: 0.0851 - val_acc: 0.9750  
Epoch 7/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0357 - acc: 0.9888 -  
val_loss: 0.0760 - val_acc: 0.9764  
Epoch 8/20  
60000/60000 [=====] - 2s 40us/step - loss: 0.0279 - acc: 0.9911 -  
val_loss: 0.0839 - val_acc: 0.9762  
Epoch 9/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0262 - acc: 0.9916 -  
val_loss: 0.0878 - val_acc: 0.9769  
Epoch 10/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0238 - acc: 0.9925 -  
val_loss: 0.0859 - val_acc: 0.9781  
Epoch 11/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0181 - acc: 0.9941 -  
val_loss: 0.0889 - val_acc: 0.9793  
Epoch 12/20  
60000/60000 [=====] - 2s 40us/step - loss: 0.0191 - acc: 0.9938 -  
val_loss: 0.0961 - val_acc: 0.9778  
Epoch 13/20  
60000/60000 [=====] - 2s 39us/step - loss: 0.0143 - acc: 0.9952 -  
val_loss: 0.1115 - val_acc: 0.9738  
Epoch 14/20  
60000/60000 [=====] - 2s 39us/step - loss: 0.0136 - acc: 0.9954 -  
val_loss: 0.1105 - val_acc: 0.9755  
Epoch 15/20  
60000/60000 [=====] - 2s 40us/step - loss: 0.0159 - acc: 0.9942 -  
val_loss: 0.1163 - val_acc: 0.9736  
Epoch 16/20  
60000/60000 [=====] - 2s 40us/step - loss: 0.0121 - acc: 0.9959 -  
val_loss: 0.1157 - val_acc: 0.9753  
Epoch 17/20  
60000/60000 [=====] - 2s 39us/step - loss: 0.0109 - acc: 0.9964 -  
val_loss: 0.1066 - val_acc: 0.9773  
Epoch 18/20  
60000/60000 [=====] - 2s 39us/step - loss: 0.0095 - acc: 0.9969 -  
val_loss: 0.0991 - val_acc: 0.9792  
Epoch 19/20  
60000/60000 [=====] - 2s 39us/step - loss: 0.0093 - acc: 0.9969 -  
val_loss: 0.1110 - val_acc: 0.9778  
Epoch 20/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0091 - acc: 0.9968 -  
val_loss: 0.1252 - val_acc: 0.9748
```


In [0]:

```
score = model3_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12519279713677378

Test accuracy: 0.9748



In [0]:

```
%matplotlib inline
w_after = model3_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + BatchNorm on hidden layers + Adam

In [0]:

```
model3_relu_batchnorm = Sequential()

model3_relu_batchnorm.add(Dense(128, activation = 'relu', input_dim = input_dim, kernel_initializer=
'glorot_normal'))
model3_relu_batchnorm.add(BatchNormalization())

model3_relu_batchnorm.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer=
'glorot_normal'))
model3_relu_batchnorm.add(BatchNormalization())

model3_relu_batchnorm.add(Dense(32, activation = 'relu', input_dim = input_dim, kernel_initializer=
'glorot_normal'))
model3_relu_batchnorm.add(BatchNormalization())

model3_relu_batchnorm.add(Dense(output_dim, activation = 'softmax'))

model3_relu_batchnorm.summary()
```

Layer (type)	Output Shape	Param #
dense_57 (Dense)	(None, 128)	100480
batch_normalization_11 (Batch Normalization)	(None, 128)	512
dense_58 (Dense)	(None, 64)	8256
batch_normalization_12 (Batch Normalization)	(None, 64)	256
dense_59 (Dense)	(None, 32)	2080
batch_normalization_13 (Batch Normalization)	(None, 32)	128

dense_60 (Dense)	(None, 10)	330
------------------	------------	-----

=====

Total params: 112,042

Trainable params: 111,594

Non-trainable params: 448

=====

In [0]:

```
model3_relu_batchnorm.compile(optimizer='adam', metrics = ['accuracy'], loss =  
'categorical_crossentropy')
```

```
history = model3_relu_batchnorm.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, v  
erbose = 1, validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 106us/step - loss: 0.3190 - acc: 0.9111 -

val_loss: 0.1437 - val_acc: 0.9570

Epoch 2/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1148 - acc: 0.9657 -

val_loss: 0.1025 - val_acc: 0.9703

Epoch 3/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0787 - acc: 0.9758 -

val_loss: 0.0985 - val_acc: 0.9709

Epoch 4/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0600 - acc: 0.9815 -

val_loss: 0.0879 - val_acc: 0.9733

Epoch 5/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0467 - acc: 0.9860 -

val_loss: 0.0918 - val_acc: 0.9732

Epoch 6/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0407 - acc: 0.9868 -

val_loss: 0.0961 - val_acc: 0.9706

Epoch 7/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0342 - acc: 0.9888 -

val_loss: 0.0818 - val_acc: 0.9760

Epoch 8/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0288 - acc: 0.9907 -

val_loss: 0.0920 - val_acc: 0.9728

Epoch 9/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0260 - acc: 0.9918 -

val_loss: 0.0791 - val_acc: 0.9773

Epoch 10/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0221 - acc: 0.9929 -

val_loss: 0.0797 - val_acc: 0.9759

Epoch 11/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0203 - acc: 0.9933 -

val_loss: 0.0832 - val_acc: 0.9761

Epoch 12/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0186 - acc: 0.9941 -

val_loss: 0.1096 - val_acc: 0.9706

Epoch 13/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0192 - acc: 0.9936 -

val_loss: 0.0861 - val_acc: 0.9760

Epoch 14/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0180 - acc: 0.9939 -

val_loss: 0.0896 - val_acc: 0.9767

Epoch 15/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0159 - acc: 0.9948 -

val_loss: 0.0955 - val_acc: 0.9753

Epoch 16/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0156 - acc: 0.9946 -

val_loss: 0.0918 - val_acc: 0.9776

Epoch 17/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0126 - acc: 0.9960 -

val_loss: 0.0896 - val_acc: 0.9773

Epoch 18/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0116 - acc: 0.9960 -

val_loss: 0.0835 - val_acc: 0.9788

Epoch 19/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0138 - acc: 0.9957 -

val_loss: 0.0803 - val_acc: 0.9778

Epoch 20/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0127 - acc: 0.9955 -
val_loss: 0.0947 - val_acc: 0.9775

In [0]:

```
score = model3_relu_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0946673099771564

Test accuracy: 0.9775



In [0]:

```
%matplotlib inline
w_after = model3_relu_batchnorm.get_weights()

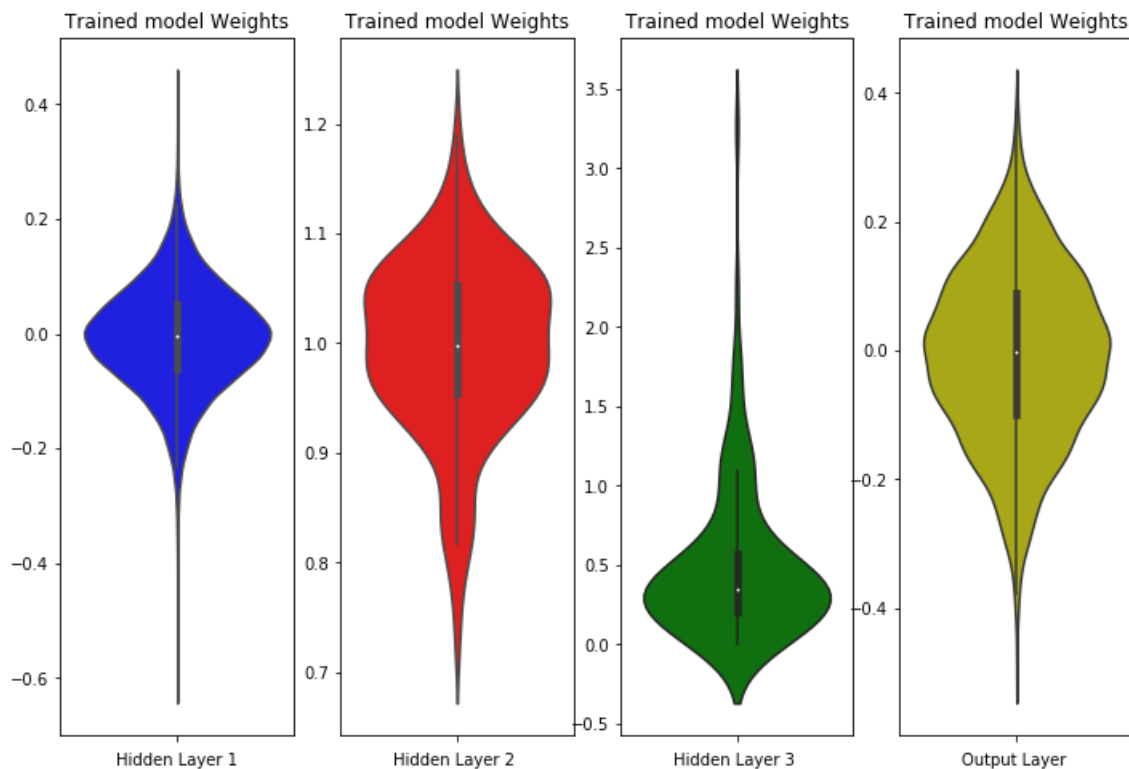
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize = (12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout + Batch Norm +Adam

In [0]:

```
model3_drop_batchnorm = Sequential()

model3_drop_batchnorm.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model3_drop_batchnorm.add(BatchNormalization())
model3_drop_batchnorm.add(Dropout(0.5))

model3_drop_batchnorm.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model3_drop_batchnorm.add(BatchNormalization())
model3_drop_batchnorm.add(Dropout(0.5))

model3_drop_batchnorm.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
model3_drop_batchnorm.add(BatchNormalization())
model3_drop_batchnorm.add(Dropout(0.5))

model3_drop_batchnorm.add(Dense(output_dim, activation='softmax'))

model3_drop_batchnorm.summary()
```

Layer (type)	Output Shape	Param #
dense_61 (Dense)	(None, 128)	100480
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_7 (Dropout)	(None, 128)	0

dense_62 (Dense)	(None, 64)	8256
batch_normalization_15 (Batch Normalization)	(None, 64)	256
dropout_8 (Dropout)	(None, 64)	0
dense_63 (Dense)	(None, 32)	2080
batch_normalization_16 (Batch Normalization)	(None, 32)	128
dropout_9 (Dropout)	(None, 32)	0
dense_64 (Dense)	(None, 10)	330
=====		
Total params: 112,042		
Trainable params: 111,594		
Non-trainable params: 448		

In [0]:

```
model3_drop_batchnorm.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model3_drop_batchnorm.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 111us/step - loss: 0.8041 - acc: 0.7578 -
val_loss: 0.2195 - val_acc: 0.9324
Epoch 2/20
60000/60000 [=====] - 4s 69us/step - loss: 0.3983 - acc: 0.8924 -
val_loss: 0.1707 - val_acc: 0.9496
Epoch 3/20
60000/60000 [=====] - 4s 69us/step - loss: 0.3370 - acc: 0.9115 -
val_loss: 0.1509 - val_acc: 0.9568
Epoch 4/20
60000/60000 [=====] - 4s 69us/step - loss: 0.3005 - acc: 0.9211 -
val_loss: 0.1393 - val_acc: 0.9594
Epoch 5/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2788 - acc: 0.9270 -
val_loss: 0.1364 - val_acc: 0.9603
Epoch 6/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2564 - acc: 0.9324 -
val_loss: 0.1178 - val_acc: 0.9664
Epoch 7/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2457 - acc: 0.9359 -
val_loss: 0.1181 - val_acc: 0.9674
Epoch 8/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2353 - acc: 0.9386 -
val_loss: 0.1151 - val_acc: 0.9688
Epoch 9/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2231 - acc: 0.9423 -
val_loss: 0.1088 - val_acc: 0.9691
Epoch 10/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2233 - acc: 0.9414 -
val_loss: 0.1122 - val_acc: 0.9682
Epoch 11/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2164 - acc: 0.9448 -
val_loss: 0.1052 - val_acc: 0.9708
Epoch 12/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2087 - acc: 0.9462 -
val_loss: 0.1076 - val_acc: 0.9706
Epoch 13/20
60000/60000 [=====] - 4s 68us/step - loss: 0.2013 - acc: 0.9479 -
val_loss: 0.1088 - val_acc: 0.9705
Epoch 14/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1961 - acc: 0.9497 -
val_loss: 0.1021 - val_acc: 0.9732
Epoch 15/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1914 - acc: 0.9504 -
val_loss: 0.1025 - val_acc: 0.9713
Epoch 16/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1886 - acc: 0.9512 -
```

```

val_loss: 0.1042 - val_acc: 0.9720
Epoch 17/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1904 - acc: 0.9505 -
val_loss: 0.0988 - val_acc: 0.9724
Epoch 18/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1861 - acc: 0.9517 -
val_loss: 0.1018 - val_acc: 0.9735
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1797 - acc: 0.9537 -
val_loss: 0.0942 - val_acc: 0.9734
Epoch 20/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1760 - acc: 0.9551 -
val_loss: 0.1012 - val_acc: 0.9715

```

In [0]:

```

score = model3_drop_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lilation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10121105482841376

Test accuracy: 0.9715



In [0]:

```

%matplotlib inline
w_after = model3_drop_batchnorm.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

```

```


fig = plt.figure(figsize= (12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



With 5 Hidden layers

Architecture used: 784 (input) - 512 - 256 - 128 - 64 - 32 - 10 (ouput)

MLP + ReLu + Adam

In [0]:

```

model5_relu = Sequential()

model5_relu.add(Dense(512, activation = 'relu', input_dim = input_dim, kernel_initializer= 'he_norm
al'))
model5_relu.add(Dense(256, activation = 'relu', input_dim = input_dim, kernel_initializer= 'he_norm
al'))
model5_relu.add(Dense(128, activation = 'relu', input_dim = input_dim, kernel_initializer= 'he_norm
al'))
model5_relu.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer= 'he_norma
l'))
model5_relu.add(Dense(32, activation = 'relu', input_dim = input_dim, kernel_initializer= 'he_norma
l'))

```



```
model5_relu.add(Dense(output_dim, activation = 'softmax'))

model5_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_65 (Dense)	(None, 512)	401920
dense_66 (Dense)	(None, 256)	131328
dense_67 (Dense)	(None, 128)	32896
dense_68 (Dense)	(None, 64)	8256
dense_69 (Dense)	(None, 32)	2080
dense_70 (Dense)	(None, 10)	330
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		

In [0]:

```
model5_relu.compile(optimizer = 'adam', metrics = ['accuracy'], loss = 'categorical_crossentropy')

history = model5_relu.fit(X_train, Y_train, batch_size= batch_size, epochs = nb_epoch, verbose = 1,
validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 5s 78us/step - loss: 0.2628 - acc: 0.9210 -
val_loss: 0.1135 - val_acc: 0.9661
Epoch 2/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0933 - acc: 0.9716 -
val_loss: 0.0857 - val_acc: 0.9739
Epoch 3/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0617 - acc: 0.9813 -
val_loss: 0.0835 - val_acc: 0.9764
Epoch 4/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0455 - acc: 0.9856 -
val_loss: 0.0929 - val_acc: 0.9723
Epoch 5/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0364 - acc: 0.9884 -
val_loss: 0.0881 - val_acc: 0.9749
Epoch 6/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0292 - acc: 0.9904 -
val_loss: 0.0666 - val_acc: 0.9797
Epoch 7/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0279 - acc: 0.9910 -
val_loss: 0.0764 - val_acc: 0.9772
Epoch 8/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0231 - acc: 0.9926 -
val_loss: 0.0841 - val_acc: 0.9779
Epoch 9/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0220 - acc: 0.9927 -
val_loss: 0.0908 - val_acc: 0.9774
Epoch 10/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0175 - acc: 0.9941 -
val_loss: 0.0791 - val_acc: 0.9807
Epoch 11/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0188 - acc: 0.9944 -
val_loss: 0.0802 - val_acc: 0.9799
Epoch 12/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0186 - acc: 0.9941 -
val_loss: 0.0721 - val_acc: 0.9817
Epoch 13/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0137 - acc: 0.9957 -
val_loss: 0.1109 - val_acc: 0.9750
Epoch 14/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0143 - acc: 0.9956 -
val_loss: 0.0976 - val_acc: 0.9790
Epoch 15/20
```

```

60000/60000 [=====] - 3s 45us/step - loss: 0.0105 - acc: 0.9967 -
val_loss: 0.0821 - val_acc: 0.9822
Epoch 16/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0137 - acc: 0.9955 -
val_loss: 0.0859 - val_acc: 0.9802
Epoch 17/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0120 - acc: 0.9963 -
val_loss: 0.0871 - val_acc: 0.9802
Epoch 18/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0105 - acc: 0.9966 -
val_loss: 0.0745 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0094 - acc: 0.9974 -
val_loss: 0.0943 - val_acc: 0.9819
Epoch 20/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0127 - acc: 0.9964 -
val_loss: 0.1076 - val_acc: 0.9777

```

In [0]:

```

score = model5_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10761469807939465
Test accuracy: 0.9777



In [0]:

```

%matplotlib inline
w_after = model5_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)

```

```

h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='purple')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='cyan')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + BatchNorm on hidden layers + Adam

In [0]:

```

model5_relu_batchnorm = Sequential()

model5_relu_batchnorm.add(Dense(512, activation = 'relu', input_dim = input_dim, kernel_initializer=
'he_uniform'))
model5_relu_batchnorm.add(BatchNormalization())

model5_relu_batchnorm.add(Dense(256, activation = 'relu', input_dim = input_dim, kernel_initializer=
'he_uniform'))
model5_relu_batchnorm.add(BatchNormalization())

```

```

model5_relu_batchnorm.add(BatchNormalization())

model5_relu_batchnorm.add(Dense(128, activation = 'relu', input_dim = input_dim, kernel_initializer=
'he_uniform'))
model5_relu_batchnorm.add(BatchNormalization())

model5_relu_batchnorm.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer=
'he_uniform'))
model5_relu_batchnorm.add(BatchNormalization())

model5_relu_batchnorm.add(Dense(32, activation = 'relu', input_dim = input_dim, kernel_initializer=
'he_uniform'))
model5_relu_batchnorm.add(BatchNormalization())

model5_relu_batchnorm.add(Dense(output_dim, activation = 'softmax'))

model5_relu_batchnorm.summary()

```

Layer (type)	Output Shape	Param #
dense_71 (Dense)	(None, 512)	401920
batch_normalization_17 (Batch Normalization)	(None, 512)	2048
dense_72 (Dense)	(None, 256)	131328
batch_normalization_18 (Batch Normalization)	(None, 256)	1024
dense_73 (Dense)	(None, 128)	32896
batch_normalization_19 (Batch Normalization)	(None, 128)	512
dense_74 (Dense)	(None, 64)	8256
batch_normalization_20 (Batch Normalization)	(None, 64)	256
dense_75 (Dense)	(None, 32)	2080
batch_normalization_21 (Batch Normalization)	(None, 32)	128
dense_76 (Dense)	(None, 10)	330
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

In [0]:

```

model5_relu_batchnorm.compile(optimizer='adam', metrics = ['accuracy'], loss =
'categorical_crossentropy')

history = model5_relu_batchnorm.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, v
erbose = 1, validation_data = (X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 141us/step - loss: 0.2496 - acc: 0.9307 -
val_loss: 0.1209 - val_acc: 0.9630
Epoch 2/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0899 - acc: 0.9733 -
val_loss: 0.0893 - val_acc: 0.9713
Epoch 3/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0632 - acc: 0.9800 -
val_loss: 0.1019 - val_acc: 0.9704
Epoch 4/20
60000/60000 [=====] - 5s 92us/step - loss: 0.0504 - acc: 0.9842 -
val_loss: 0.0886 - val_acc: 0.9735
Epoch 5/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0402 - acc: 0.9878 -

```

```

val_loss: 0.0748 - val_acc: 0.9773
Epoch 6/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0346 - acc: 0.9883 -
val_loss: 0.0799 - val_acc: 0.9766
Epoch 7/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0301 - acc: 0.9900 -
val_loss: 0.0952 - val_acc: 0.9731
Epoch 8/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0289 - acc: 0.9905 -
val_loss: 0.0750 - val_acc: 0.9791
Epoch 9/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0254 - acc: 0.9920 -
val_loss: 0.0821 - val_acc: 0.9775
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0229 - acc: 0.9932 -
val_loss: 0.0902 - val_acc: 0.9749
Epoch 11/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0197 - acc: 0.9936 -
val_loss: 0.0827 - val_acc: 0.9805
Epoch 12/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0180 - acc: 0.9940 -
val_loss: 0.0966 - val_acc: 0.9747
Epoch 13/20
60000/60000 [=====] - 5s 92us/step - loss: 0.0215 - acc: 0.9931 -
val_loss: 0.0757 - val_acc: 0.9812
Epoch 14/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0164 - acc: 0.9944 -
val_loss: 0.0729 - val_acc: 0.9807
Epoch 15/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0176 - acc: 0.9940 -
val_loss: 0.0747 - val_acc: 0.9817
Epoch 16/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0135 - acc: 0.9953 -
val_loss: 0.0915 - val_acc: 0.9787
Epoch 17/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0159 - acc: 0.9946 -
val_loss: 0.0810 - val_acc: 0.9803
Epoch 18/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0121 - acc: 0.9958 -
val_loss: 0.0768 - val_acc: 0.9810
Epoch 19/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0117 - acc: 0.9963 -
val_loss: 0.0740 - val_acc: 0.9815
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0115 - acc: 0.9961 -
val_loss: 0.0821 - val_acc: 0.9810

```

In [0]:

```

score = model5_relu_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08212392491589125
Test accuracy: 0.981



In [0]:

```
%matplotlib inline
w_after = model5_relu_batchnorm.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='purple')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='cyan')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + Dropout + Batch Norm +Adam

In [0]:

```
model5_drop_batchnorm = Sequential()

model5_drop_batchnorm.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(128, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(64, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(32, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(output_dim, activation='softmax'))

model5_drop_batchnorm.summary()
```

Layer (type)	Output Shape	Param #
dense_77 (Dense)	(None, 512)	401920
batch_normalization_22 (Batch Normalization)	(None, 512)	2048
dropout_10 (Dropout)	(None, 512)	0
dense_78 (Dense)	(None, 256)	131328
batch_normalization_23 (Batch Normalization)	(None, 256)	1024
dropout_11 (Dropout)	(None, 256)	0
dense_79 (Dense)	(None, 128)	32896
batch_normalization_24 (Batch Normalization)	(None, 128)	512
dropout_12 (Dropout)	(None, 128)	0
dense_80 (Dense)	(None, 64)	8256
batch_normalization_25 (Batch Normalization)	(None, 64)	256
dropout_13 (Dropout)	(None, 64)	0
dense_81 (Dense)	(None, 32)	2080
batch_normalization_26 (Batch Normalization)	(None, 32)	128

dropout_14 (Dropout)	(None, 32)	0
dense_82 (Dense)	(None, 10)	330
=====		
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		
=====		

In [0]:

```
model5_drop_batchnorm.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model5_drop_batchnorm.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 148us/step - loss: 1.4684 - acc: 0.5186 -
val_loss: 0.3500 - val_acc: 0.9124
Epoch 2/20
60000/60000 [=====] - 5s 91us/step - loss: 0.6071 - acc: 0.8267 -
val_loss: 0.1998 - val_acc: 0.9451
Epoch 3/20
60000/60000 [=====] - 6s 92us/step - loss: 0.4087 - acc: 0.8955 -
val_loss: 0.1712 - val_acc: 0.9553
Epoch 4/20
60000/60000 [=====] - 6s 92us/step - loss: 0.3354 - acc: 0.9188 -
val_loss: 0.1467 - val_acc: 0.9635
Epoch 5/20
60000/60000 [=====] - 6s 92us/step - loss: 0.2902 - acc: 0.9311 -
val_loss: 0.1325 - val_acc: 0.9671
Epoch 6/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2676 - acc: 0.9374 -
val_loss: 0.1239 - val_acc: 0.9691
Epoch 7/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2399 - acc: 0.9430 -
val_loss: 0.1128 - val_acc: 0.9720
Epoch 8/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2205 - acc: 0.9496 -
val_loss: 0.1050 - val_acc: 0.9735
Epoch 9/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2037 - acc: 0.9531 -
val_loss: 0.1067 - val_acc: 0.9737
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1914 - acc: 0.9560 -
val_loss: 0.0976 - val_acc: 0.9774
Epoch 11/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1806 - acc: 0.9585 -
val_loss: 0.0867 - val_acc: 0.9781
Epoch 12/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1772 - acc: 0.9602 -
val_loss: 0.0888 - val_acc: 0.9781
Epoch 13/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1635 - acc: 0.9637 -
val_loss: 0.0875 - val_acc: 0.9799
Epoch 14/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1598 - acc: 0.9634 -
val_loss: 0.0852 - val_acc: 0.9789
Epoch 15/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1533 - acc: 0.9658 -
val_loss: 0.0832 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1423 - acc: 0.9678 -
val_loss: 0.0834 - val_acc: 0.9797
Epoch 17/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1411 - acc: 0.9689 -
val_loss: 0.0837 - val_acc: 0.9807
Epoch 18/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1394 - acc: 0.9690 -
val_loss: 0.0854 - val_acc: 0.9805
Epoch 19/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1385 - acc: 0.9687 -
val_loss: 0.0796 - val_acc: 0.9806
```


Epoch 20/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1336 - acc: 0.9707 -
val_loss: 0.0829 - val_acc: 0.9800

In [0]:

```
score = model5_drop_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08293465451768134
Test accuracy: 0.98



In [0]:

```
%matplotlib inline
w_after = model5_drop_batchnorm.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
```


```
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='purple')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='cyan')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLu + RMSprop

In [0]:

```
model2_relu = Sequential()

model2_relu.add(Dense(256, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_uniform'))
model2_relu.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_uniform'))
model2_relu.add(Dense(output_dim, activation = 'softmax'))

model2_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_83 (Dense)	(None, 256)	200960
dense_84 (Dense)	(None, 64)	16448
dense_85 (Dense)	(None, 10)	650
Total params: 218,058		
Trainable params: 218,058		
Non-trainable params: 0		

In [0]:

```
model2_relu.compile(optimizer = 'RMSprop', metrics = ['accuracy'], loss =  
'categorical_crossentropy')  
  
history = model2_relu.fit(X_train, Y_train, batch_size= batch_size, epochs = nb_epoch, verbose = 1,  
validation_data = (X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples  
Epoch 1/20  
60000/60000 [=====] - 5s 76us/step - loss: 0.2732 - acc: 0.9207 -  
val_loss: 0.1480 - val_acc: 0.9533  
Epoch 2/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.1059 - acc: 0.9677 -  
val_loss: 0.0837 - val_acc: 0.9752  
Epoch 3/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0710 - acc: 0.9784 -  
val_loss: 0.0902 - val_acc: 0.9718  
Epoch 4/20  
60000/60000 [=====] - 2s 36us/step - loss: 0.0524 - acc: 0.9839 -  
val_loss: 0.0674 - val_acc: 0.9795  
Epoch 5/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0405 - acc: 0.9877 -  
val_loss: 0.0686 - val_acc: 0.9796  
Epoch 6/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0314 - acc: 0.9904 -  
val_loss: 0.0749 - val_acc: 0.9791  
Epoch 7/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0252 - acc: 0.9919 -  
val_loss: 0.0726 - val_acc: 0.9797  
Epoch 8/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0197 - acc: 0.9940 -  
val_loss: 0.0660 - val_acc: 0.9829  
Epoch 9/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0156 - acc: 0.9949 -  
val_loss: 0.0771 - val_acc: 0.9818  
Epoch 10/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0131 - acc: 0.9960 -  
val_loss: 0.0868 - val_acc: 0.9803  
Epoch 11/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0110 - acc: 0.9966 -  
val_loss: 0.0794 - val_acc: 0.9827  
Epoch 12/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0087 - acc: 0.9973 -  
val_loss: 0.0782 - val_acc: 0.9819  
Epoch 13/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0076 - acc: 0.9977 -  
val_loss: 0.0851 - val_acc: 0.9805  
Epoch 14/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0060 - acc: 0.9979 -  
val_loss: 0.0904 - val_acc: 0.9815  
Epoch 15/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0053 - acc: 0.9983 -  
val_loss: 0.0951 - val_acc: 0.9808  
Epoch 16/20  
60000/60000 [=====] - 2s 36us/step - loss: 0.0042 - acc: 0.9987 -  
val_loss: 0.1038 - val_acc: 0.9804  
Epoch 17/20  
60000/60000 [=====] - 2s 35us/step - loss: 0.0045 - acc: 0.9984 -  
val_loss: 0.1029 - val_acc: 0.9820  
Epoch 18/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0034 - acc: 0.9990 -  
val_loss: 0.1090 - val_acc: 0.9808  
Epoch 19/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0029 - acc: 0.9991 -  
val_loss: 0.0974 - val_acc: 0.9824  
Epoch 20/20  
60000/60000 [=====] - 2s 34us/step - loss: 0.0028 - acc: 0.9992 -  
val_loss: 0.0999 - val_acc: 0.9833
```

In [0]:

```
score = model2_relu.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])
```

```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0998641132327826

Test accuracy: 0.9833



In [0]:

```

%matplotlib inline
w_after = model2_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

MLP + ReLu + Adadelata

In [0]:

```
model3_relu = Sequential()

model3_relu.add(Dense(128, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_normal'))
model3_relu.add(Dense(64, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_normal'))
model3_relu.add(Dense(32, activation = 'relu', input_dim = input_dim, kernel_initializer= 'glorot_normal'))
model3_relu.add(Dense(output_dim, activation = 'softmax'))

model3_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_86 (Dense)	(None, 128)	100480
dense_87 (Dense)	(None, 64)	8256
dense_88 (Dense)	(None, 32)	2080
dense_89 (Dense)	(None, 10)	330
Total params: 111,146		
Trainable params: 111,146		
Non-trainable params: 0		

In [0]:

```
model3_relu.compile(optimizer = 'Adadelata', metrics = ['accuracy'], loss = 'categorical_crossentropy')

history = model3_relu.fit(X_train, Y_train, batch_size= batch_size, epochs = nb_epoch, verbose = 1, validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 5s 83us/step - loss: 0.3476 - acc: 0.8968 - val_loss: 0.1580 - val_acc: 0.9529
Epoch 2/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1320 - acc: 0.9609 - val_loss: 0.1127 - val_acc: 0.9636
Epoch 3/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0906 - acc: 0.9723 - val_loss: 0.0882 - val_acc: 0.9725
Epoch 4/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0702 - acc: 0.9783 - val_loss: 0.0878 - val_acc: 0.9737
Epoch 5/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0550 - acc: 0.9834 - val_loss: 0.0809 - val_acc: 0.9742
Epoch 6/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0442 - acc: 0.9864 - val_loss: 0.0777 - val_acc: 0.9772
Epoch 7/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0369 - acc: 0.9885 - val_loss: 0.0752 - val_acc: 0.9767
Epoch 8/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0297 - acc: 0.9915 - val_loss: 0.0788 - val_acc: 0.9774
Epoch 9/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0251 - acc: 0.9928 - val_loss: 0.0721 - val_acc: 0.9779
Epoch 10/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0206 - acc: 0.9937 - val_loss: 0.0731 - val_acc: 0.9775
Epoch 11/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0168 - acc: 0.9952 - val_loss: 0.0755 - val_acc: 0.9771
Epoch 12/20
```

```

Epoch 12/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0136 - acc: 0.9959 -
val_loss: 0.0775 - val_acc: 0.9786
Epoch 13/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0117 - acc: 0.9966 -
val_loss: 0.0767 - val_acc: 0.9790
Epoch 14/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0092 - acc: 0.9974 -
val_loss: 0.0797 - val_acc: 0.9785
Epoch 15/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0070 - acc: 0.9983 -
val_loss: 0.0825 - val_acc: 0.9801
Epoch 16/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0062 - acc: 0.9983 -
val_loss: 0.0856 - val_acc: 0.9781
Epoch 17/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0051 - acc: 0.9987 -
val_loss: 0.0847 - val_acc: 0.9793
Epoch 18/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0036 - acc: 0.9992 -
val_loss: 0.0899 - val_acc: 0.9799
Epoch 19/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0028 - acc: 0.9994 -
val_loss: 0.0841 - val_acc: 0.9806
Epoch 20/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0023 - acc: 0.9996 -
val_loss: 0.0918 - val_acc: 0.9791

```

In [0]:

```

score = model3_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09183809580091215

Test accuracy: 0.9791



2.5 5.0 7.5 10.0 12.5 15.0 17.5 20.0
epoch

In [0]:

```
%matplotlib inline
w_after = model3_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(12,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout + Batch Norm + Nadam

In [0]:

```
model5_drop_batchnorm = Sequential()

model5_drop_batchnorm.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initialize
r='he_normal'))
```

```

# he_normal //
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(128, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(64, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(32, activation='relu', kernel_initializer='he_normal'))
model5_drop_batchnorm.add(BatchNormalization())
model5_drop_batchnorm.add(Dropout(0.5))

model5_drop_batchnorm.add(Dense(output_dim, activation='softmax'))

model5_drop_batchnorm.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_90 (Dense)	(None, 512)	401920
batch_normalization_27 (Batch Normalization)	(None, 512)	2048
dropout_15 (Dropout)	(None, 512)	0
dense_91 (Dense)	(None, 256)	131328
batch_normalization_28 (Batch Normalization)	(None, 256)	1024
dropout_16 (Dropout)	(None, 256)	0
dense_92 (Dense)	(None, 128)	32896
batch_normalization_29 (Batch Normalization)	(None, 128)	512
dropout_17 (Dropout)	(None, 128)	0
dense_93 (Dense)	(None, 64)	8256
batch_normalization_30 (Batch Normalization)	(None, 64)	256
dropout_18 (Dropout)	(None, 64)	0
dense_94 (Dense)	(None, 32)	2080
batch_normalization_31 (Batch Normalization)	(None, 32)	128
dropout_19 (Dropout)	(None, 32)	0
dense_95 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

In [0]:

```

model5_drop_batchnorm.compile(optimizer='Nadam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model5_drop_batchnorm.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 165us/step - loss: 1.1588 - acc: 0.6263 - val_loss: 1.1588 - val_acc: 0.6263

```



```

oss: 0.2429 - val_acc: 0.9330
Epoch 2/20
60000/60000 [=====] - 6s 99us/step - loss: 0.4527 - acc: 0.8792 -
val_loss: 0.1712 - val_acc: 0.9536
Epoch 3/20
60000/60000 [=====] - 6s 100us/step - loss: 0.3354 - acc: 0.9165 -
val_loss: 0.1522 - val_acc: 0.9614
Epoch 4/20
60000/60000 [=====] - 6s 99us/step - loss: 0.2882 - acc: 0.9307 -
val_loss: 0.1282 - val_acc: 0.9681
Epoch 5/20
60000/60000 [=====] - 6s 99us/step - loss: 0.2559 - acc: 0.9388 -
val_loss: 0.1190 - val_acc: 0.9695
Epoch 6/20
60000/60000 [=====] - 6s 100us/step - loss: 0.2396 - acc: 0.9440 -
val_loss: 0.1110 - val_acc: 0.9722
Epoch 7/20
60000/60000 [=====] - 6s 100us/step - loss: 0.2223 - acc: 0.9491 -
val_loss: 0.1076 - val_acc: 0.9728
Epoch 8/20
60000/60000 [=====] - 6s 100us/step - loss: 0.2062 - acc: 0.9520 -
val_loss: 0.0915 - val_acc: 0.9773
Epoch 9/20
60000/60000 [=====] - 6s 99us/step - loss: 0.1964 - acc: 0.9535 -
val_loss: 0.0991 - val_acc: 0.9764
Epoch 10/20
60000/60000 [=====] - 6s 102us/step - loss: 0.1910 - acc: 0.9563 -
val_loss: 0.0998 - val_acc: 0.9762
Epoch 11/20
60000/60000 [=====] - 6s 102us/step - loss: 0.1822 - acc: 0.9580 -
val_loss: 0.0976 - val_acc: 0.9763
Epoch 12/20
60000/60000 [=====] - 6s 101us/step - loss: 0.1761 - acc: 0.9604 -
val_loss: 0.1017 - val_acc: 0.9757
Epoch 13/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1674 - acc: 0.9618 -
val_loss: 0.0943 - val_acc: 0.9779
Epoch 14/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1635 - acc: 0.9628 -
val_loss: 0.0880 - val_acc: 0.9788
Epoch 15/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1574 - acc: 0.9642 -
val_loss: 0.0919 - val_acc: 0.9793
Epoch 16/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1518 - acc: 0.9645 -
val_loss: 0.0866 - val_acc: 0.9791
Epoch 17/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1520 - acc: 0.9664 -
val_loss: 0.0812 - val_acc: 0.9808
Epoch 18/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1449 - acc: 0.9671 -
val_loss: 0.0888 - val_acc: 0.9801
Epoch 19/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1405 - acc: 0.9681 -
val_loss: 0.0797 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 6s 100us/step - loss: 0.1364 - acc: 0.9680 -
val_loss: 0.0895 - val_acc: 0.9794

```

In [0]:

```

score = model5_drop_batchnorm.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

```


```
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08949376999950036

Test accuracy: 0.9794



In [0]:

```
%matplotlib inline
w_after = model5_drop_batchnorm.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,8))
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='purple')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='cyan')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ')
plt.show()
```



Summary:

1. Experimented different MLP architectures on the MNIST dataset:

- 2 hidden layers - 784 (input) - 256 - 64 - 10 (ouput)
- 3 hidden layers - 784 (input) - 128 - 64 - 32- 10 (output)
- 5 hidden layers - 784 (input) - 512 - 256 - 128 - 64 - 32 - 10 (output)

1. Initialized differrent weight vectors using:

- glorot-normal
- glorot-uniform
- he-normal
- he-uniform
- random_normal
- random_uniform

2. For every architecture, plotted epoch vs loss for training and validation data.

3. For sanity check, plotted violin plots of weights after training the model.

4. Also, performed batch normalization and dropout and it resulted in increase in the accuracy.

5. Conducted a comparison to see what performs better: normbatch before dropout vs dropout before normbatch. Found out dropout before normbatch performed slightly better by looking at the test accuracy.

In [4]:

```
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["No. of layers", "Architecture", "MLP type", "Optimizer", "Test Score", "Test Accuracy"]

x.add_row([2, '784-256-64-10', "MLP + ReLu + Adam", "Adam", 0.09701326997896285, 0.9795])
x.add_row([2, '784-256-64-10', "MLP + BatchNorm on hidden layers + Adam", "Adam", 0.0820598181706071, 0.9799])
x.add_row([2, '784-256-64-10', "MLP + Batch Norm + Dropout + Adam", "Adam", 0.0703403844089189, 0.9796])
x.add_row([2, '784-256-64-10', "MLP + Dropout + Batch Norm + Adam", "Adam", 0.06499798040131573, 0.9817])

x.add_row([3, '784-128-64-32-10', "MLP + ReLu + Adam", "Adam", 0.10691093345177942, 0.9765])
x.add_row([3, '784-128-64-32-10', "MLP + BatchNorm on hidden layers + Adam", "Adam", 0.11158925297728202, 0.9762])
x.add_row([3, '784-128-64-32-10', "MLP + Dropout + Batch Norm + Adam", "Adam", 0.09398230195709038, 0.9758])

x.add_row([5, '784-512-256-128-64-32-10', "MLP + ReLu + Adam", "Adam", 0.09024974172847265, 0.9824])
x.add_row([5, '784-512-256-128-64-32-10', "MLP + BatchNorm on hidden layers + Adam", "Adam",
```

```

0.08782999148442759, 0.9791])
x.add_row([5, '784-512-256-128-64-32-10',"MLP + Dropout + Batch Norm +Adam", "Adam" ,
0.07287385137048549, 0.9845])

# models with different optimizers
x.add_row([2, '784-256-64-10',"MLP + ReLu + RMSprop", "RMSprop",0.0998641132327826, 0.9833])
x.add_row([3, '784-128-64-32-10',"MLP + ReLu + Adadelata", "Adadelata", 0.09183809580091215,
0.9791])
x.add_row([5, '784-512-256-128-64-32-10',"MLP + Dropout + Batch Norm + Nadam", "Nadam" ,
0.08949376999950036, 0.9794])

print(x)

```

No. of layers		Architecture		MLP type		Optimizer	
Test Score		Test Accuracy					
0.09701326997896285	2	0.9795	784-256-64-10	MLP + ReLu + Adam		Adam	
0.0820598181706071	2	0.9799	784-256-64-10	MLP + BatchNorm on hidden layers + Adam		Adam	
0.0703403844089189	2	0.9796	784-256-64-10	MLP + Batch Norm + Dropout + Adam		Adam	
0.06499798040131573	2	0.9817	784-256-64-10	MLP + Dropout + Batch Norm +Adam		Adam	
0.10691093345177942	3	0.9765	784-128-64-32-10	MLP + ReLu + Adam		Adam	
0.11158925297728202	3	0.9762	784-128-64-32-10	MLP + BatchNorm on hidden layers + Adam		Adam	
0.09398230195709038	3	0.9758	784-128-64-32-10	MLP + Dropout + Batch Norm +Adam		Adam	
0.09024974172847265	5	0.9824	784-512-256-128-64-32-10	MLP + ReLu + Adam		Adam	
0.08782999148442759	5	0.9791	784-512-256-128-64-32-10	MLP + BatchNorm on hidden layers + Adam		Adam	
0.07287385137048549	5	0.9845	784-512-256-128-64-32-10	MLP + Dropout + Batch Norm +Adam		Adam	
0.0998641132327826	2	0.9833	784-256-64-10	MLP + ReLu + RMSprop		RMSprop	
0.09183809580091215	3	0.9791	784-128-64-32-10	MLP + ReLu + Adadelata		Adadelata	
0.08949376999950036	5	0.9794	784-512-256-128-64-32-10	MLP + Dropout + Batch Norm + Nadam		Nadam	